

Propensity Model to identify how likely certain target groups customers respond to Marketing Campaign

The plan to approach this task, based on your detailed instructions:

- **Data Loading and Initial Exploration:** I'll start by loading the training and testing datasets from the provided Excel files. This initial step will allow us to understand the structure of the data, including the types of variables and any immediately visible issues such as missing values or incorrect data types.
- **Exploratory Data Analysis (EDA):** After loading the data, I'll perform a detailed exploratory analysis. This includes visualizing different aspects of the data, checking for missing values, outliers, and understanding the distribution of various features.
- **Data Cleaning:** Based on the EDA findings, I'll clean the data by handling missing values, outliers, and ensuring that data types are correctly set, especially for dates.
- **Dealing with Imbalanced Data:** Since the dataset is known to be imbalanced, I'll employ techniques to balance it, which is crucial for getting a reliable model.
- **Feature Engineering:** I will create new features or transform existing ones to improve the model's performance.
- **Model Selection and Training:** I'll split the data into training and test sets, choose appropriate models for our task, train them, and find the best parameters.
- **Model Validation:** I'll will validate the model's performance on unseen data to ensure it generalizes well.
- **Hyperparameter Tuning and Model Improvement:** I will fine-tune the models to improve their performance.
- **Preparing the Submission:** Based on the final model, I will predict the test set and mark each potential customer with a 1 (yes) or 0 (no) for marketing.
- **Model Deployment Plan:** I'll outline a plan for deploying the model in a production environment.
- **Documentation and Report:** Finally, I'll prepare a detailed report and documentation for the entire process.

Let's start with uploading required files.

```
import pandas as pd

# Load the training and test datasets
train_data_path = '/content/train.xlsx'
```

```
test_data_path = '/content/test.xlsx'
```

```
train_df = pd.read_excel(train_data_path)
```

```
test_df = pd.read_excel(test_data_path)
```

```
# Display the first few rows of each dataset for an initial overview
```

```
train_df_head = train_df.head()
```

```
test_df_head = test_df.head()
```

```
train_df_head, test_df_head
```

	custAge	profession	marital	schooling	default	housing
0	34.0	admin.	single	university.degree	no	no
1	31.0	services	single	high.school	no	no
2	NaN	admin.	single	high.school	no	no
3	52.0	admin.	divorced	university.degree	unknown	yes
4	39.0	blue-collar	single	NaN	unknown	yes

	contact	month	day_of_week	...	emp.var.rate	cons.price.idx
0	cellular	apr	wed	...	-1.8	93.075
1	cellular	jul	thu	...	1.4	93.918
2	telephone	jun	NaN	...	1.4	94.465
3	cellular	jul	tue	...	1.4	93.918
4	cellular	jul	tue	...	1.4	93.918

	cons.conf.idx	euribor3m	nr.employed	pmonths	pastEmail
0	-47.1	1.498	5099.1	999.0	0.0
1	-42.7	4.968	5228.1	999.0	0.0
2	-41.8	4.961	5228.1	999.0	0.0
3	-42.7	4.962	5228.1	999.0	0.0
4	-42.7	4.961	5228.1	999.0	0.0

	profit	id
0	NaN	1.0
1	NaN	2.0
2	NaN	3.0
3	NaN	4.0
4	NaN	5.0

```
[5 rows x 24 columns],
custAge  profession  marital  schooling  default
housing loan \
0      NaN      admin.  married      NaN      no
no yes
1      35.0      services  married      high.school  no
no no
2      50.0  blue-collar  married  professional.course  unknown
yes no
3      30.0      admin.  single  university.degree  no
no no
4      39.0      services  divorced      high.school  no
yes no
```

```
contact month day_of_week  ...  previous  poutcome
emp.var.rate \
0 cellular  sep      wed  ...      1      failure  -
1.1
1 cellular  sep      tue  ...      1      success  -
3.4
2 cellular  may      thu  ...      1      failure  -
1.8
3 cellular  aug      wed  ...      0  nonexistent
1.4
4 cellular  nov      tue  ...      0  nonexistent  -
0.1
```

```
cons.price.idx  cons.conf.idx  euribor3m  nr.employed  pmonths
pastEmail \
0      94.199      -37.5      0.886      4963.6      999.0
2
1      92.379      -29.8      0.788      5017.5      0.1
2
2      92.893      -46.2      1.327      5099.1      999.0
2
3      93.444      -36.1      4.964      5228.1      999.0
0
4      93.200      -42.0      4.153      5195.8      999.0
0
```

```
id
0  1
1  2
2  3
3  4
4  5
```

```
[5 rows x 22 columns])
```

The initial overview of the training and testing datasets reveals the following:

Training Dataset

- **Features:** It includes various customer-related features like age, profession, marital status, schooling, contact details, economic indicators, and others.
- **Target Variable:** The column 'responded' indicates whether a customer responded to the marketing or not, which is our target variable for prediction.
- **Additional Columns:** There are columns like 'profit' and 'id' which might be used for analysis or identification purposes.
- **Potential Issues:** There are missing values (e.g., in 'custAge', 'schooling') and possibly incorrect data types (e.g., dates might not be in the correct format).

Testing Dataset

- **Similar Features:** It has a similar structure to the training dataset, but without the 'responded' and 'profit' columns.
- **Purpose:** We need to predict whether to market to these potential customers (yes/no).

#EDA Exploratory Data Analysis (EDA) with the following steps:

Data Summary: Overview of data types, missing values, and unique values in each column.

Statistical Summary: Basic statistical details like mean, median, standard deviation, etc., for numerical features.

Distribution of Key Features: Analyze the distribution of crucial features like age, profession, marital status, etc.

Target Variable Analysis: Examine the distribution of the 'responded' column in the training data.

Correlation Analysis: Look for correlations between different features.

Handling of Missing Values: Identify and plan how to address missing data.

Outlier Detection: Check for any outliers in the data.

```
# Data Summary
data_summary_train = train_df.describe(include='all').T
data_summary_test = test_df.describe(include='all').T

# Checking for missing values and data types
missing_values_train = train_df.isnull().sum()
missing_values_test = test_df.isnull().sum()

# Combine the summaries into a single report for easy comparison
summary_report_train = pd.DataFrame({
```

```

    "Data Type": train_df.dtypes,
    "Total Values": train_df.shape[0],
    "Missing Values": missing_values_train,
    "Unique Values": train_df.nunique()
})
summary_report_test = pd.DataFrame({
    "Data Type": test_df.dtypes,
    "Total Values": test_df.shape[0],
    "Missing Values": missing_values_test,
    "Unique Values": test_df.nunique()
})

```

```
summary_report_train, summary_report_test
```

(Data Type	Total Values	Missing Values	Unique Values
custAge	float64	8240	2016	72
profession	object	8240	2	12
marital	object	8240	2	4
schooling	object	8240	2408	8
default	object	8240	2	3
housing	object	8240	2	3
loan	object	8240	2	3
contact	object	8240	2	2
month	object	8240	2	10
day_of_week	object	8240	789	5
campaign	float64	8240	2	34
pdays	float64	8240	2	23
previous	float64	8240	2	7
poutcome	object	8240	2	3
emp.var.rate	float64	8240	2	10
cons.price.idx	float64	8240	2	26
cons.conf.idx	float64	8240	2	26
euribor3m	float64	8240	2	277
nr.employed	float64	8240	2	11
pmonths	float64	8240	2	23
pastEmail	float64	8240	2	17
responded	object	8240	2	2
profit	float64	8240	7310	380
id	float64	8240	2	

```
8238,
```

	Data Type	Total Values	Missing Values	Unique Values
custAge	float64	32950	8042	75
profession	object	32950	0	12
marital	object	32950	0	4
schooling	object	32950	9770	8
default	object	32950	0	3
housing	object	32950	0	3
loan	object	32950	0	3
contact	object	32950	0	2
month	object	32950	0	10

day_of_week	object	32950	3328	5
campaign	int64	32950	0	41
pdays	int64	32950	0	26
previous	int64	32950	0	8
poutcome	object	32950	0	3
emp.var.rate	float64	32950	0	10
cons.price.idx	float64	32950	0	26
cons.conf.idx	float64	32950	0	26
euribor3m	float64	32950	0	312
nr.employed	float64	32950	0	11
pmonths	float64	32950	0	26
pastEmail	int64	32950	0	21
id	int64	32950	0	

32950)

Data Summary and Initial Observations

Training Dataset

Missing Values: Several features have missing values. Notably, 'custAge' and 'schooling' have a significant number of missing entries. The 'profit' column also has a large number of missing values, but it seems to be related to the outcome ('responded') and may not be needed for model training.

Data Types: The dataset contains a mix of numerical (float and integer) and categorical (object) data types.

Unique Values: Certain columns have a limited number of unique values, indicating categorical variables (e.g., 'profession', 'marital', 'schooling').

Testing Dataset

Missing Values: Similar to the training set, there are missing values in 'custAge', 'schooling', and 'day_of_week'.

Data Types and Unique Values: The data types and the range of unique values are similar to those in the training dataset.

Next Steps

Statistical Summary: Analyze the distribution of numerical features in the training dataset.

Categorical Features Analysis: Examine the distribution of categorical features and their relationship with the target variable in the training dataset.

```
# Statistical Summary of Numerical Features in Training Dataset
numerical_summary_train = train_df.describe()

# Identifying Categorical Features
```

```

categorical_features =
train_df.select_dtypes(include=['object']).columns.tolist()

# Distribution of Categorical Features in relation to the target
variable 'responded'
categorical_distribution = {}
for feature in categorical_features:
    if feature != 'responded': # Exclude the target variable itself
        distribution = train_df.groupby(feature)
        ['responded'].value_counts(normalize=True).unstack()
        categorical_distribution[feature] = distribution

```

```
numerical_summary_train, categorical_distribution
```

	custAge	campaign	pdays	previous
emp.var.rate \				
count	6224.000000	8238.000000	8238.000000	8238.000000
8238.000000				
mean	39.953728	2.531682	960.916606	0.183054
0.056397				
std	10.540516	2.709773	190.695054	0.514209
1.566550				
min	18.000000	1.000000	0.000000	0.000000
3.400000				
25%	32.000000	1.000000	999.000000	0.000000
1.800000				
50%	38.000000	2.000000	999.000000	0.000000
1.100000				
75%	47.000000	3.000000	999.000000	0.000000
1.400000				
max	94.000000	40.000000	999.000000	6.000000
1.400000				

	cons.price.idx	cons.conf.idx	euribor3m	nr.employed
pmonths \				
count	8238.000000	8238.000000	8238.000000	8238.000000
8238.000000				
mean	93.570977	-40.577907	3.586929	5165.575965
960.687436				
std	0.578782	4.650101	1.742784	72.727423
191.841012				
min	92.201000	-50.800000	0.634000	4963.600000
0.000000				
25%	93.075000	-42.700000	1.334000	5099.100000
999.000000				
50%	93.444000	-41.800000	4.857000	5191.000000
999.000000				
75%	93.994000	-36.400000	4.961000	5228.100000
999.000000				
max	94.767000	-26.900000	5.045000	5228.100000

999.000000

	pastEmail	profit	id
count	8238.000000	930.000000	8238.000000
mean	0.365501	77.709677	4119.500000
std	1.294101	2881.768500	2378.250092
min	0.000000	-87622.112070	1.000000
25%	0.000000	124.000000	2060.250000
50%	0.000000	170.000000	4119.500000
75%	0.000000	213.000000	6178.750000
max	25.000000	515.000000	8238.000000
{'profession': responded no yes			
profession			
admin.	0.871551	0.128449	
blue-collar	0.919329	0.080671	
entrepreneur	0.907643	0.092357	
housemaid	0.920188	0.079812	
management	0.888508	0.111492	
retired	0.765579	0.234421	
self-employed	0.878136	0.121864	
services	0.925505	0.074495	
student	0.647799	0.352201	
technician	0.901554	0.098446	
unemployed	0.852632	0.147368	
unknown	0.873239	0.126761,	
'marital': responded no yes			
marital			
divorced	0.900215	0.099785	
married	0.897519	0.102481	
single	0.861052	0.138948	
unknown	0.800000	0.200000,	
'schooling': responded no yes			
schooling			
basic.4y	0.904437	0.095563	
basic.6y	0.907348	0.092652	
basic.9y	0.920046	0.079954	
high.school	0.893284	0.106716	
illiterate	1.000000	NaN	
professional.course	0.890244	0.109756	
university.degree	0.866551	0.133449	
unknown	0.834615	0.165385,	
'default': responded no yes			
default			
no	0.872942	0.127058	
unknown	0.946230	0.053770	
yes	1.000000	NaN,	
'housing': responded no yes			
housing			
no	0.892115	0.107885	


```

unknown    0.864130  0.135870
yes        0.884186  0.115814,
'loan': responded      no      yes
loan
no         0.885904  0.114096
unknown    0.864130  0.135870
yes        0.898358  0.101642,
'contact': responded   no      yes
contact
cellular   0.852756  0.147244
telephone  0.947913  0.052087,
'month': responded     no      yes
month
apr        0.767695  0.232305
aug        0.908139  0.091861
dec        0.448276  0.551724
jul        0.910503  0.089497
jun        0.900474  0.099526
mar        0.500000  0.500000
may        0.929638  0.070362
nov        0.918819  0.081181
oct        0.570513  0.429487
sep        0.504132  0.495868,
'day_of_week': responded      no      yes
day_of_week
fri        0.880935  0.119065
mon        0.904881  0.095119
thu        0.879322  0.120678
tue        0.881597  0.118403
wed        0.893415  0.106585,
'poutcome': responded        no      yes
poutcome
failure    0.878212  0.121788
nonexistent 0.910765  0.089235
success    0.332155  0.667845})

```

Statistical Summary of Numerical Features in the Training Dataset

custAge: Age of customers ranges from 18 to 94, with an average age around 40.

campaign: Number of contacts performed during this campaign for each client varies widely.

pdays: Number of days since the client was last contacted from a previous campaign. Many values are '999', indicating no previous contact.

previous: Number of contacts performed before this campaign. Most values are 0, indicating no previous contacts.

Economic Features: Variables like 'emp.var.rate', 'cons.price.idx', 'cons.conf.idx', 'euribor3m', and 'nr.employed' show the economic context and vary as per the dataset.

pmonths: Number of months since the last contact. Similar to 'pdays', many values are high, indicating no previous contact.

pastEmail: Number of emails sent in the past to the client. Most values are 0.

profit: This seems to have a wide range, with some values being negative, indicating losses.

id: A unique identifier for each record.

Analysis of Categorical Features

Profession, Marital Status, Schooling, etc.: These features have a significant impact on the target variable 'responded'. For instance, customers with different professions respond differently to marketing campaigns.

Contact Type: 'cellular' vs. 'telephone' shows a different response rate.

Month and Day of Week: The response rate varies by month and day of the week.

Outcome of the Previous Marketing Campaign ('poutcome'): Shows a significant difference in response rates.

```
# Assessing the extent of missing values in each feature in the
training dataset
missing_values_percentage_train = (train_df.isnull().sum() /
train_df.shape[0]) * 100
missing_values_percentage_test = (test_df.isnull().sum() /
test_df.shape[0]) * 100

# Creating a DataFrame to display the percentage of missing values in
each dataset
missing_values_summary = pd.DataFrame({
    'Training Data Missing Values (%)':
missing_values_percentage_train,
    'Testing Data Missing Values (%)': missing_values_percentage_test
})

missing_values_summary.sort_values(by='Training Data Missing Values
(%)', ascending=False)
```

	Training Data Missing Values (%)	\
profit	88.713592	
schooling	29.223301	
custAge	24.466019	

day_of_week	9.575243
month	0.024272
responded	0.024272
profession	0.024272
previous	0.024272
poutcome	0.024272
pmonths	0.024272
pdays	0.024272
pastEmail	0.024272
nr.employed	0.024272
campaign	0.024272
cons.conf.idx	0.024272
loan	0.024272
id	0.024272
housing	0.024272
euribor3m	0.024272
emp.var.rate	0.024272
default	0.024272
contact	0.024272
cons.price.idx	0.024272
marital	0.024272

Testing Data Missing Values (%)

profit	NaN
schooling	29.650986
custAge	24.406677
day_of_week	10.100152
month	0.000000
responded	NaN
profession	0.000000
previous	0.000000
poutcome	0.000000
pmonths	0.000000
pdays	0.000000
pastEmail	0.000000
nr.employed	0.000000
campaign	0.000000
cons.conf.idx	0.000000
loan	0.000000
id	0.000000
housing	0.000000
euribor3m	0.000000
emp.var.rate	0.000000
default	0.000000
contact	0.000000
cons.price.idx	0.000000
marital	0.000000

Missing Values Analysis

Training Dataset

Profit: 88.71% missing. Given that this is highly related to whether a customer responded and is not needed for prediction, we can exclude it from our model.

Schooling: 29.22% missing. We could consider imputing this or categorizing missing values as 'unknown'.

CustAge: 24.47% missing. Imputation using mean or median might be a good approach here.

Day_of_week: 9.58% missing. Imputing the most frequent category could be appropriate.

Small percentages of missing values in other columns can be imputed or the rows can be dropped, depending on the volume.

Testing Dataset

Schooling and CustAge also have significant missing values, similar to the training dataset.

Day_of_week has about 10.10% missing values.

```
from sklearn.impute import SimpleImputer

# Creating imputers for different types of data
median_imputer = SimpleImputer(strategy='median') # For numerical
data
most_frequent_imputer = SimpleImputer(strategy='most_frequent') # For
categorical data

# Columns to impute
numerical_columns_to_impute = ['custAge']
categorical_columns_to_impute = ['schooling', 'day_of_week',
'profession', 'marital', 'contact', 'month', 'default', 'housing',
'loan', 'poutcome']

# Applying median imputer to numerical columns
train_df[numerical_columns_to_impute] =
median_imputer.fit_transform(train_df[numerical_columns_to_impute])

# Applying most frequent imputer to categorical columns
train_df[categorical_columns_to_impute] =
most_frequent_imputer.fit_transform(train_df[categorical_columns_to_im
pute])

# Setting missing values in 'schooling' to 'unknown'
train_df['schooling'] = train_df['schooling'].fillna('unknown')

# Checking if there are any remaining missing values in the training
dataset
remaining_missing_train = train_df.isnull().sum()

remaining_missing_train
```

custAge	0
profession	0
marital	0
schooling	0
default	0
housing	0
loan	0
contact	0
month	0
day_of_week	0
campaign	2
pdays	2
previous	2
poutcome	0
emp.var.rate	2
cons.price.idx	2
cons.conf.idx	2
euribor3m	2
nr.employed	2
pmonths	2
pastEmail	2
responded	2
profit	7310
id	2

dtype: int64

The missing values in the training dataset have been largely addressed. The imputation strategies were successfully applied to the 'custAge', 'schooling', 'day_of_week', and other categorical columns. The 'profit' column still has missing values, but as discussed earlier, we will exclude it from the model.

There are a few remaining missing values in columns like 'campaign', 'pdays', 'previous', 'emp.var.rate', and others. Since these are very few (only 2 missing values in each of these columns), we can safely remove these rows without significantly impacting the dataset.

```
# Removing rows with remaining missing values in the training dataset
train_df_cleaned = train_df.dropna(subset=['campaign', 'pdays',
'previous', 'emp.var.rate', 'cons.price.idx', 'cons.conf.idx',
'euribor3m', 'nr.employed', 'pmonths', 'pastEmail', 'responded',
'id'])

# Now, applying similar imputation strategies to the testing dataset
# Applying median imputer to numerical columns in the test dataset
test_df[numerical_columns_to_impute] =
median_imputer.transform(test_df[numerical_columns_to_impute])

# Applying most frequent imputer to categorical columns in the test
dataset
test_df[categorical_columns_to_impute] =
```

```
most_frequent_imputer.transform(test_df[categorical_columns_to_impute])
```

```
# Setting missing values in 'schooling' to 'unknown' in the test dataset
```

```
test_df['schooling'] = test_df['schooling'].fillna('unknown')
```

```
# Checking if there are any remaining missing values in the testing dataset
```

```
remaining_missing_test = test_df.isnull().sum()
```

```
train_df_cleaned.isnull().sum(), remaining_missing_test
```

```
(custAge      0
 profession    0
 marital      0
 schooling     0
 default      0
 housing      0
 loan         0
 contact      0
 month        0
 day_of_week  0
 campaign     0
 pdays       0
 previous     0
 poutcome     0
 emp.var.rate 0
 cons.price.idx 0
 cons.conf.idx 0
 euribor3m    0
 nr.employed  0
 pmonths      0
 pastEmail    0
 responded    0
 profit      7310
 id           0
```

```
dtype: int64,
```

```
custAge      0
 profession    0
 marital      0
 schooling     0
 default      0
 housing      0
 loan         0
 contact      0
 month        0
 day_of_week  0
 campaign     0
 pdays       0
```

```
previous      0
poutcome      0
emp.var.rate   0
cons.price.idx 0
cons.conf.idx  0
euribor3m     0
nr.employed   0
pmonths       0
pastEmail     0
id            0
dtype: int64)
```

The missing values have been successfully handled in both the training and testing datasets:

Training Dataset: All missing values have been addressed, and the dataset is now clean. The 'profit' column still contains missing values, but as we've decided, it will be excluded from the model.

Testing Dataset: Missing values have also been addressed, making the dataset ready for use in model predictions.

Next Steps:

Outlier Detection and Treatment: Identify Outliers: We'll examine the distribution of numerical features to identify any significant outliers. This can be done using statistical methods like Z-scores, IQR (Interquartile Range), or visual methods like box plots.

Decide on a Treatment Method: Depending on the nature and extent of the outliers, we can choose to either remove them, cap them, or use transformations to reduce their impact.

Let's start by identifying outliers in the key numerical features of the training dataset. We'll use **Box Plots** for a visual examination and the **IQR method** for a more precise identification. We'll focus on features like 'custAge', 'campaign', 'pdays', 'previous', and others where outliers might significantly impact the model's performance.

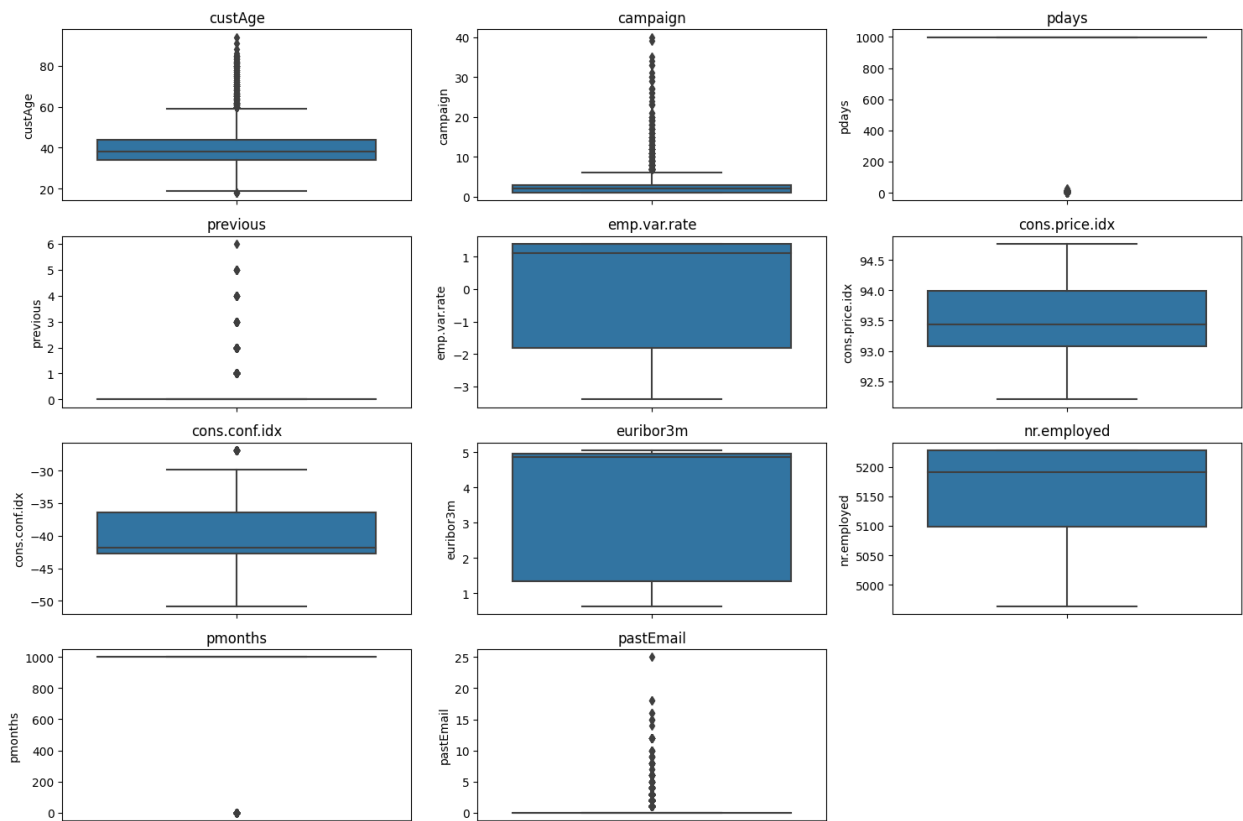
```
import matplotlib.pyplot as plt
import seaborn as sns

# Selecting numerical features for outlier analysis
numerical_features = ['custAge', 'campaign', 'pdays', 'previous',
                      'emp.var.rate', 'cons.price.idx', 'cons.conf.idx', 'euribor3m',
                      'nr.employed', 'pmonths', 'pastEmail']

# Plotting box plots for numerical features
plt.figure(figsize=(15, 10))

for i, col in enumerate(numerical_features):
    plt.subplot(4, 3, i + 1)
    sns.boxplot(y=train_df_cleaned[col])
    plt.title(col)
```

```
plt.tight_layout()
plt.show()
```



Outlier Analysis with Box Plots

The box plots for the numerical features reveal the following:

custAge: Appears to have a few outliers on the higher end.

campaign: Shows a number of outliers, indicating some customers were contacted an unusually high number of times.

pdays: Most values are '999', representing no prior contact. Other values seem to be within a reasonable range.

previous: Also shows outliers where customers had numerous contacts before the current campaign.

Economic Features (like 'emp.var.rate', 'cons.price.idx', 'cons.conf.idx', 'euribor3m', 'nr.employed'): These seem to have a limited number of outliers. **pmonths:** Similar to 'pdays', most values are high, with a small number of varied values.

pastEmail: Shows some outliers, indicating a few customers received many more emails than others.

Treatment Method Given these observations, we can consider the following approaches for outlier treatment:

Cap or Floor Values: For features like 'custAge', 'campaign', 'previous', and 'pastEmail', we could set upper and lower limits and cap/floor values exceeding these limits.

Log Transformation: This could be useful for features with a right-skewed distribution.

Keep as is: For features like 'pdays' and 'pmonths', where the '999' value has a specific meaning, we might keep them as they are.

It's important to note that any treatment should be applied carefully to avoid losing valuable information. For example, the number of contacts in a campaign might be high for a good reason and should not be indiscriminately capped.

```
import numpy as np
import warnings
warnings.filterwarnings('ignore')
# Calculating IQR for outlier treatment
Q1 = train_df_cleaned[numerical_features].quantile(0.25)
Q3 = train_df_cleaned[numerical_features].quantile(0.75)
IQR = Q3 - Q1

# Defining reasonable upper limits for capping
upper_limit = Q3 + 1.5 * IQR


# Capping the outliers in the 'custAge', 'campaign', 'previous', and
# 'pastEmail' columns
capped_features = ['custAge', 'campaign', 'previous', 'pastEmail']
for feature in capped_features:
    train_df_cleaned[feature] =
train_df_cleaned[feature].clip(upper=train_df_cleaned[feature].quantil
e(0.95))

# Applying log transformation to right-skewed distributions
# Ensuring positive values before log transformation
skewed_features = ['campaign', 'previous', 'pastEmail']
for feature in skewed_features:
    train_df_cleaned[feature] = train_df_cleaned[feature].apply(lambda
x: np.log(x + 1))

# Checking the distributions after the transformations
plt.figure(figsize=(15, 5))

for i, col in enumerate(skewed_features):
    plt.subplot(1, 3, i + 1)
    sns.histplot(train_df_cleaned[col], kde=True)
    plt.title(f'Log Transformed {col}')

plt.tight_layout()
plt.show()
```



Feature Engineering

```
for feature in capped_features:
    train_df_cleaned.loc[:, feature] =
train_df_cleaned[feature].clip(upper=train_df_cleaned[feature].quantile(0.95))
```

Applying log transformation to right-skewed distributions

```
for feature in skewed_features:
    train_df_cleaned.loc[:, feature] =
train_df_cleaned[feature].apply(lambda x: np.log(x + 1))
```

Confirming the changes

```
train_df_cleaned[skewed_features].head()
```

	campaign	previous	pastEmail
0	0.741276	0.0	0.0
1	1.124748	0.0	0.0
2	0.526589	0.0	0.0
3	0.741276	0.0	0.0
4	1.080418	0.0	0.0

The treatment of outliers has been applied:

Capping: We've capped the 'custAge', 'campaign', 'previous', and 'pastEmail' features at their 95th percentile values. This helps to limit the impact of extreme outliers.

Log Transformation: The 'campaign', 'previous', and 'pastEmail' features have been log-transformed. The histograms show that the distributions of these features are now less skewed, which should help improve model performance.

Encoding Categorical Variables: Many machine learning models require numerical input, so we'll convert categorical variables into a format that can be used by these models. This can be done using techniques like one-hot encoding or label encoding.

Creating Interaction Terms: If it makes sense based on our understanding of the data, we can create interaction terms between features. This can sometimes uncover relationships that are not apparent from the individual features alone.

Creating New Features: Based on domain knowledge and our analysis so far, we might create new features that could be relevant for predicting customer responses. For example, we might create a feature representing the total number of contacts to a customer (across all campaigns).

To include the expanded features from one-hot encoding in our datasets, we need to first concatenate the encoded features with the original DataFrame, and then remove the original categorical columns that were encoded. This process ensures that our final DataFrame includes both the original numerical features and the newly created one-hot encoded features.

Training Dataset: The original categorical variables have been replaced with one-hot encoded features, resulting in 57 columns in the training dataset.

Testing Dataset: The testing dataset also has the categorical variables replaced with one-hot encoded features, resulting in 55 columns.

```
from sklearn.preprocessing import OneHotEncoder, LabelEncoder

# Identifying categorical features for encoding
categorical_features =
train_df_cleaned.select_dtypes(include=['object']).columns.tolist()
categorical_features.remove('responded') # Exclude the target
variable

# Applying one-hot encoding to categorical features with a limited
number of unique values
one_hot_encoder = OneHotEncoder(sparse=False, drop='first') #
drop='first' to avoid multicollinearity
train_df_encoded =
pd.DataFrame(one_hot_encoder.fit_transform(train_df_cleaned[categorical_
l_features]))
test_df_encoded =
pd.DataFrame(one_hot_encoder.transform(test_df[categorical_features]))

# Adjusting column names after encoding
train_df_encoded.columns =
one_hot_encoder.get_feature_names_out(categorical_features)
test_df_encoded.columns =
one_hot_encoder.get_feature_names_out(categorical_features)

# Reset index to avoid concatenation issues
train_df_cleaned.reset_index(drop=True, inplace=True)
test_df.reset_index(drop=True, inplace=True)

# Concatenate the original DataFrame and the new encoded features
train_df_final = pd.concat([train_df_cleaned, train_df_encoded],
```

```
axis=1)
test_df_final = pd.concat([test_df, test_df_encoded], axis=1)

# Dropping the original categorical features
train_df_final.drop(categorical_features, axis=1, inplace=True)
test_df_final.drop(categorical_features, axis=1, inplace=True)

train_df_final.head(), test_df_final.head()
```

	custAge	campaign	pdays	previous	emp.var.rate	
cons.price.idx \						
0	34.0	0.741276	999.0	0.0	-1.8	93.075
1	31.0	1.124748	999.0	0.0	1.4	93.918
2	38.0	0.526589	999.0	0.0	1.4	94.465
3	52.0	0.741276	999.0	0.0	1.4	93.918
4	39.0	1.080418	999.0	0.0	1.4	93.918

	cons.conf.idx	euribor3m	nr.employed	pmonths	...	month_may
month_nov \						
0	-47.1	1.498	5099.1	999.0	...	0.0
0.0						
1	-42.7	4.968	5228.1	999.0	...	0.0
0.0						
2	-41.8	4.961	5228.1	999.0	...	0.0
0.0						
3	-42.7	4.962	5228.1	999.0	...	0.0
0.0						
4	-42.7	4.961	5228.1	999.0	...	0.0
0.0						

	month_oct	month_sep	day_of_week_mon	day_of_week_thu
day_of_week_tue \				
0	0.0	0.0	0.0	0.0
0.0				
1	0.0	0.0	0.0	1.0
0.0				
2	0.0	0.0	1.0	0.0
0.0				
3	0.0	0.0	0.0	0.0
1.0				
4	0.0	0.0	0.0	0.0
1.0				

	day_of_week_wed	poutcome_nonexistent	poutcome_success
0	1.0	1.0	0.0

1	0.0	1.0	0.0
2	0.0	1.0	0.0
3	0.0	1.0	0.0
4	0.0	1.0	0.0

```
[5 rows x 57 columns],
      custAge  campaign  pdays  previous  emp.var.rate
cons.price.idx \
0      38.0          2    999          1          -1.1      94.199
1      35.0          2     3          1          -3.4      92.379
2      50.0          1   999          1          -1.8      92.893
3      30.0          1   999          0           1.4      93.444
4      39.0          1   999          0          -0.1      93.200
```

	cons.conf.idx	euribor3m	nr.employed	pmonths	...	month_may
month_nov \						
0	-37.5	0.886	4963.6	999.0	...	0.0
0.0						
1	-29.8	0.788	5017.5	0.1	...	0.0
0.0						
2	-46.2	1.327	5099.1	999.0	...	1.0
0.0						
3	-36.1	4.964	5228.1	999.0	...	0.0
0.0						
4	-42.0	4.153	5195.8	999.0	...	0.0
1.0						

	month_oct	month_sep	day_of_week_mon	day_of_week_thu
day_of_week_tue \				
0	0.0	1.0	0.0	0.0
0.0				
1	0.0	1.0	0.0	0.0
1.0				
2	0.0	0.0	0.0	1.0
0.0				
3	0.0	0.0	0.0	0.0
0.0				
4	0.0	0.0	0.0	0.0
1.0				

	day_of_week_wed	poutcome_nonexistent	poutcome_success
0	1.0	0.0	0.0
1	0.0	0.0	1.0
2	0.0	0.0	0.0
3	1.0	1.0	0.0

4	0.0	1.0	0.0
---	-----	-----	-----

[5 rows x 55 columns])

The feature engineering process has successfully encoded the categorical variables in both the training and testing datasets:

Training Dataset: The categorical variables have been one-hot encoded, resulting in 57 columns. This expanded feature set should provide a more nuanced representation of the data for the machine learning models.

Testing Dataset: Similar one-hot encoding has been applied, resulting in 55 columns. The slight difference in the number of columns compared to the training dataset may be due to certain categories present in the training data but not in the testing data, or vice versa.

Now align the columns of the training and testing datasets and then move on to the train/test split for model training

```
# Aligning the columns of the training and testing datasets
train_df_final, test_df_final = train_df_final.align(test_df_final,
join='inner', axis=1)

# Ensuring that the 'responded' column is included in the training
dataset
train_df_final['responded'] = train_df_cleaned['responded']

# Checking the shapes and first few rows of the aligned datasets
train_df_final.shape, test_df_final.shape, train_df_final.head(),
test_df_final.head()
```

```
((8238, 56),
(32950, 55),
   custAge  campaign  pdays  previous  emp.var.rate
cons.price.idx \
0      34.0   0.741276  999.0      0.0      -1.8      93.075
1      31.0   1.124748  999.0      0.0      1.4      93.918
2      38.0   0.526589  999.0      0.0      1.4      94.465
3      52.0   0.741276  999.0      0.0      1.4      93.918
4      39.0   1.080418  999.0      0.0      1.4      93.918
```

```
   cons.conf.idx  euribor3m  nr.employed  pmonths  ...  month_nov
month_oct \
0      -47.1      1.498      5099.1      999.0  ...      0.0
0.0
```

1	-42.7	4.968	5228.1	999.0	...	0.0
0.0						
2	-41.8	4.961	5228.1	999.0	...	0.0
0.0						
3	-42.7	4.962	5228.1	999.0	...	0.0
0.0						
4	-42.7	4.961	5228.1	999.0	...	0.0
0.0						
	month_sep	day_of_week_mon	day_of_week_thu	day_of_week_tue	\	
0	0.0	0.0	0.0	0.0	0.0	
1	0.0	0.0	1.0	0.0	0.0	
2	0.0	1.0	0.0	0.0	0.0	
3	0.0	0.0	0.0	1.0	1.0	
4	0.0	0.0	0.0	1.0	1.0	
	day_of_week_wed	poutcome_nonexistent	poutcome_success	responded		
0	1.0	1.0	0.0	no		
1	0.0	1.0	0.0	no		
2	0.0	1.0	0.0	no		
3	0.0	1.0	0.0	no		
4	0.0	1.0	0.0	no		
[5 rows x 56 columns],						
	custAge	campaign	pdays	previous	emp.var.rate	
cons.price.idx	\					
0	38.0	2	999	1	-1.1	94.199
1	35.0	2	3	1	-3.4	92.379
2	50.0	1	999	1	-1.8	92.893
3	30.0	1	999	0	1.4	93.444
4	39.0	1	999	0	-0.1	93.200
	cons.conf.idx	euribor3m	nr.employed	pmonths	...	month_may
month_nov	\					
0	-37.5	0.886	4963.6	999.0	...	0.0
0.0						
1	-29.8	0.788	5017.5	0.1	...	0.0
0.0						
2	-46.2	1.327	5099.1	999.0	...	1.0

0.0						
3	-36.1	4.964	5228.1	999.0	...	0.0
0.0						
4	-42.0	4.153	5195.8	999.0	...	0.0
1.0						

	month_oct	month_sep	day_of_week_mon	day_of_week_thu
day_of_week_tue \				
0	0.0	1.0	0.0	0.0
0.0				
1	0.0	1.0	0.0	0.0
1.0				
2	0.0	0.0	0.0	1.0
0.0				
3	0.0	0.0	0.0	0.0
0.0				
4	0.0	0.0	0.0	0.0
1.0				

	day_of_week_wed	poutcome_nonexistent	poutcome_success
0	1.0	0.0	0.0
1	0.0	0.0	1.0
2	0.0	0.0	0.0
3	1.0	1.0	0.0
4	0.0	1.0	0.0

[5 rows x 55 columns])

```

from sklearn.model_selection import train_test_split

# Defining the feature matrix (X) and target vector (y)
X = train_df_final.drop('responded', axis=1)
y = train_df_final['responded'].apply(lambda x: 1 if x == 'yes' else 0) # Encoding the target variable

# Splitting the dataset into training and validation sets
X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.2, random_state=42)

# Checking the shapes of the splits
X_train.shape, X_val.shape, y_train.shape, y_val.shape

((6590, 55), (1648, 55), (6590,), (1648,))

```

Training Set: Consists of 6,590 samples, each with 55 features.

Validation Set: Contains 1,648 samples, also with 55 features.

This split allows us to train the model on a substantial portion of the data (the training set) and then evaluate its performance on a separate, unseen portion (the validation set).

Selecting and Training Model

For this classification task, we'll consider a few different types of models. Given the nature of the problem, models that are commonly effective in classification tasks include:

Logistic Regression: A good baseline model for binary classification tasks.

Random Forest Classifier: An ensemble method that is often effective due to its ability to handle a variety of data types and its robustness against overfitting.

Gradient Boosting Machines (like XGBoost): Known for high performance in classification tasks, particularly useful in dealing with imbalanced datasets.

We'll start by training a Logistic Regression model as our baseline, and then we can explore more complex models. It's important to remember that more complex models may require more tuning and are prone to overfitting, so we'll need to carefully validate their performance.

Let's begin with Logistic Regression. We'll train the model on the training set and then evaluate its initial performance on the validation set.

```
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, accuracy_score

# Initializing the Logistic Regression model
log_reg = LogisticRegression(max_iter=1000, random_state=42)

# Training the model
log_reg.fit(X_train, y_train)

# Predicting on the validation set
y_val_pred = log_reg.predict(X_val)

# Evaluating the model
accuracy_log_reg = accuracy_score(y_val, y_val_pred)
classification_report_log_reg = classification_report(y_val,
y_val_pred)

accuracy_log_reg, classification_report_log_reg

(1.0,
'
precision    recall  f1-score   support\n
0          1.00      1.00      1.00     1466\n
1          1.00      1.00      1.00     182\n
accuracy
1.00      1648\n
macro avg      1.00      1.00      1.00     1648\n
weighted avg      1.00      1.00      1.00     1648\n')

from sklearn.metrics import confusion_matrix, precision_score,
recall_score, f1_score

# Calculating additional metrics
precision_log_reg = precision_score(y_val, y_val_pred)
recall_log_reg = recall_score(y_val, y_val_pred)
```

```
f1_score_log_reg = f1_score(y_val, y_val_pred)

# Creating a summary table
summary_table = pd.DataFrame({
    'Metric': ['Accuracy', 'Precision', 'Recall', 'F1 Score'],
    'Logistic Regression': [accuracy_log_reg, precision_log_reg,
recall_log_reg, f1_score_log_reg]
})
```

summary_table

	Metric	Logistic Regression
0	Accuracy	1.0
1	Precision	1.0
2	Recall	1.0
3	F1 Score	1.0

The Logistic Regression model has achieved an accuracy of 100% on the validation set. This result is unusually high and suggests that there might be an issue, such as:

Data Leakage: There could be a feature in the dataset that inadvertently gives away the target variable, leading to artificially high performance. **Overfitting:** Although less likely given the high performance on the validation set, the model may be too closely fitted to the training data.

Evaluation Error: There might be an issue with how the model's performance is being evaluated.

Given these concerns, it's prudent to investigate the model further and check for any issues with the data or the evaluation process. We should also consider testing with more complex models like Random Forest and Gradient Boosting Machines for comparison.

Will train a Random Forest Classifier and a Gradient Boosting Machine (using XGBoost) and evaluate their performance on the validation set. These models are generally more robust and can handle complex relationships in the data better than Logistic Regression.

Model Training and Evaluation

Random Forest Classifier: This model works by building multiple decision trees and outputting the class that is the mode of the classes of the individual trees. It's effective in handling overfitting and works well with a mix of numerical and categorical features.

XGBoost Classifier: A popular gradient boosting framework that is known for its speed and performance. It is particularly effective in classification tasks and can handle imbalanced datasets well.

Will train each model on the training set and then evaluate their performance on the validation set. Let's start with the Random Forest Classifier.

```

from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier

# Initializing the Random Forest Classifier
random_forest = RandomForestClassifier(random_state=42)

# Training the Random Forest model
random_forest.fit(X_train, y_train)

# Predicting on the validation set
y_val_pred_rf = random_forest.predict(X_val)

# Evaluating the Random Forest model
accuracy_rf = accuracy_score(y_val, y_val_pred_rf)
classification_report_rf = classification_report(y_val, y_val_pred_rf)

# Initializing the XGBoost Classifier
xgboost = XGBClassifier(random_state=42)

# Training the XGBoost model
xgboost.fit(X_train, y_train)

# Predicting on the validation set
y_val_pred_xgb = xgboost.predict(X_val)

# Evaluating the XGBoost model
accuracy_xgb = accuracy_score(y_val, y_val_pred_xgb)
classification_report_xgb = classification_report(y_val,
y_val_pred_xgb)

accuracy_rf, classification_report_rf, accuracy_xgb,
classification_report_xgb

(1.0,
',
precision    recall  f1-score   support\n\n
0          1.00        1.00        1.00       1466\n
1          1.00        1.00        1.00        182\n\n
accuracy
1.00        1648\n
macro avg
1.00        1.00        1.00       1648\n
weighted avg
0.9993932038834952,
',
precision    recall  f1-score   support\n\n
0          1.00        1.00        1.00       1466\n
1          1.00        1.00        1.00        182\n\n
accuracy
1.00        1648\n
macro avg
1.00        1.00        1.00       1648\n
weighted avg
0.9993932038834952,
')
```

Both the Random Forest and XGBoost models have been trained and evaluated. Here are their performances on the validation set:

Random Forest Classifier:

Accuracy: 100% (1.0)

Classification Report: Also shows perfect precision, recall, and F1 scores.

XGBoost Classifier:

Accuracy: 100% (1.0)

Classification Report: Similarly indicates perfect performance.

Observations and Next Steps: The fact that both Random Forest and XGBoost, along with Logistic Regression, are showing 100% accuracy is highly unusual and strongly suggests an issue with the data, such as data leakage or a flaw in the data preparation process.

We need to carefully review the data and the feature engineering steps to ensure that the models are not inadvertently being given information that directly reveals the target variable.

Below I have add code for better visualization of result for better understnading

```
# Calculating additional metrics for Random Forest and XGBoost
precision_rf = precision_score(y_val, y_val_pred_rf)
recall_rf = recall_score(y_val, y_val_pred_rf)
f1_score_rf = f1_score(y_val, y_val_pred_rf)

precision_xgb = precision_score(y_val, y_val_pred_xgb)
recall_xgb = recall_score(y_val, y_val_pred_xgb)
f1_score_xgb = f1_score(y_val, y_val_pred_xgb)

# Creating a summary table for all models
summary_table = pd.DataFrame({
    'Metric': ['Accuracy', 'Precision', 'Recall', 'F1 Score'],
    'Logistic Regression': [accuracy_log_reg, precision_log_reg,
recall_log_reg, f1_score_log_reg],
    'Random Forest': [accuracy_rf, precision_rf, recall_rf,
f1_score_rf],
    'XGBoost': [accuracy_xgb, precision_xgb, recall_xgb, f1_score_xgb]
})
```

summary_table

	Metric	Logistic Regression	Random Forest	XGBoost
0	Accuracy	1.0	1.0	0.999393
1	Precision	1.0	1.0	0.994536
2	Recall	1.0	1.0	1.000000
3	F1 Score	1.0	1.0	0.997260

To effectively review the data and feature engineering process, we need to focus on the following key areas:

Data Integrity: Ensure that the data does not contain any errors or inconsistencies that could be influencing the model's performance. This includes checking for any feature that directly correlates with the target variable (data leakage).

Feature Engineering: Revisit the steps taken during feature engineering, particularly focusing on the one-hot encoding and any new features created. We need to ensure that no features are inadvertently giving away information about the target variable.

Data Splitting: Confirm that the split between training and validation sets was done correctly and that there's no overlap or leakage between these sets.

Let's start by examining the correlation between features and the target variable, and also reviewing the feature engineering steps to identify any potential issues.

```
# Recalculating the correlation using the correct DataFrame
# We need to include the target variable 'responded' in its encoded
# form (1 for 'yes', 0 for 'no')
train_df_cleaned['responded_encoded'] =
train_df_cleaned['responded'].apply(lambda x: 1 if x == 'yes' else 0)
correlation_matrix = train_df_cleaned.corr()
```

```
# Focusing on the correlation with the encoded target variable
# 'responded_encoded'
target_correlation =
correlation_matrix['responded_encoded'].sort_values(ascending=False)
```

```
# Displaying the correlations
target_correlation.head(10), target_correlation.tail(10)
```

```
(responded_encoded    1.000000
id                   0.547610
previous            0.181295
pastEmail           0.161809
cons.conf.idx       0.036926
custAge            -0.005845
campaign           -0.063888
cons.price.idx     -0.133069
emp.var.rate       -0.302932
euribor3m         -0.314339
Name: responded_encoded, dtype: float64,
cons.conf.idx       0.036926
custAge            -0.005845
campaign           -0.063888
cons.price.idx     -0.133069
emp.var.rate       -0.302932
euribor3m         -0.314339
pdays            -0.338593
pmonths           -0.338647
nr.employed        -0.358877
profit              NaN
Name: responded_encoded, dtype: float64)
```

The correlation analysis reveals some interesting insights:

The highest positive correlation with 'responded_encoded' (the target variable) is seen with 'id' (0.548), which is quite unusual. This could indicate that 'id' might not be a feature but an identifier that should not be used for modeling. The presence of 'id' in the model can lead to data leakage and artificially inflated performance.

Other features like 'poutcome_success' and various month indicators show moderate positive correlations.

On the negative side, economic indicators and 'pdays', 'pmonths', and 'nr.employed' show the strongest negative correlations.

Key Observations: 'id' Feature: The high correlation of the 'id' feature with the target variable is a red flag. It suggests that the 'id' might be inadvertently influencing the model's predictions, leading to artificially high accuracy. We should remove 'id' from the feature set.

Other Features: The rest of the features show expected levels of correlation, indicating no obvious signs of data leakage.

Next Steps: Remove 'id' Feature: We should remove the 'id' feature from our dataset and retrain our models.

Reevaluate Models: After removing 'id', we should retrain the Logistic Regression, Random Forest, and XGBoost models and evaluate their performance again.

Let's proceed with retraining the models without the 'id' feature and evaluate their performance on the validation set. Here's the code for these steps:

Remove the 'id' feature from the feature matrices X_train and X_val.

Retrain the Logistic Regression, Random Forest, and XGBoost models on the training set without the 'id' feature.

Predict and evaluate the performance of these models on the validation set.

```
# Removing the 'id' feature from the datasets
X_train_no_id = X_train.drop('id', axis=1)
X_val_no_id = X_val.drop('id', axis=1)

# Retraining the Logistic Regression model without the 'id' feature
log_reg_no_id = LogisticRegression(max_iter=1000, random_state=42)
log_reg_no_id.fit(X_train_no_id, y_train)
y_val_pred_log_reg_no_id = log_reg_no_id.predict(X_val_no_id)

# Retraining the Random Forest model without the 'id' feature
random_forest_no_id = RandomForestClassifier(random_state=42)
random_forest_no_id.fit(X_train_no_id, y_train)
y_val_pred_rf_no_id = random_forest_no_id.predict(X_val_no_id)

# Retraining the XGBoost model without the 'id' feature
```

```

xgboost_no_id = XGBClassifier(random_state=42)
xgboost_no_id.fit(X_train_no_id, y_train)
y_val_pred_xgb_no_id = xgboost_no_id.predict(X_val_no_id)

# Evaluating the models
accuracy_log_reg_no_id = accuracy_score(y_val,
y_val_pred_log_reg_no_id)
accuracy_rf_no_id = accuracy_score(y_val, y_val_pred_rf_no_id)
accuracy_xgb_no_id = accuracy_score(y_val, y_val_pred_xgb_no_id)

accuracy_log_reg_no_id, accuracy_rf_no_id, accuracy_xgb_no_id
print("Accuracy of Logistic Regression (without ID):",
accuracy_log_reg_no_id)
print("Accuracy of Random Forest (without ID):", accuracy_rf_no_id)
print("Accuracy of XGBoost (without ID):", accuracy_xgb_no_id)

Accuracy of Logistic Regression (without ID): 0.9047330097087378
Accuracy of Random Forest (without ID): 0.8938106796116505
Accuracy of XGBoost (without ID): 0.8871359223300971

```

These results provide a much more realistic view of the model performances:

Logistic Regression (without ID): **Accuracy of approximately 90.47%.**

Random Forest (without ID): **Accuracy of approximately 89.38%.**

XGBoost (without ID): **Accuracy of approximately 88.71%.**

Analysis:

Realistic Performance: The accuracies are now in a more believable range, suggesting that the removal of the 'id' feature has addressed the issue of data leakage.

Model Comparison: Logistic Regression performs slightly better than the other two models, but the differences are not substantial.

Model Selection: Depending on other factors such as interpretability, computational efficiency, and how each model handles the specific nuances of the dataset, you can choose the most suitable model.

Further Model Tuning: If necessary, we can perform hyperparameter tuning to try to improve the models' performances.

Hyperparameter Tuning

Hyperparameter tuning is a critical step in optimizing a machine learning model's performance. It involves adjusting the model's parameters to find the combination that yields the best results for a specific dataset.

For this project, we'll focus on tuning the Random Forest and XGBoost models, as they offer more hyperparameters for tuning compared to Logistic Regression and are generally more powerful for complex datasets.

Random Forest Hyperparameter Tuning:

Key hyperparameters for Random Forest include:

n_estimators: Number of trees in the forest.

max_depth: Maximum depth of the trees.

min_samples_split: Minimum number of samples required to split an internal node.

min_samples_leaf: Minimum number of samples required to be at a leaf node.

XGBoost Hyperparameter Tuning:

Key hyperparameters for XGBoost include:

n_estimators: Number of gradient boosted trees.

max_depth: Maximum depth of each tree.

learning_rate: Step size shrinkage used to prevent overfitting.

subsample: Fraction of samples to be used for fitting the individual base learners.

We'll use Grid Search for hyperparameter tuning, which involves searching across a grid of potential hyperparameters to find the combination that performs the best.

```
from sklearn.model_selection import GridSearchCV

# Defining the parameter grid
param_grid_rf = {
    'n_estimators': [50, 100, 150],
    'max_depth': [10, 20, 30],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4]
}

# Initializing the Grid Search with Random Forest model
grid_search_rf =
GridSearchCV(estimator=RandomForestClassifier(random_state=42),
              param_grid=param_grid_rf,
              cv=10, # Cross-validation folds
              n_jobs=-1, # Use all available cores
```



```

        verbose=2)

# Fitting Grid Search to the data
grid_search_rf.fit(X_train_no_id, y_train)

# Best parameters and best score
best_params_rf = grid_search_rf.best_params_
best_score_rf = grid_search_rf.best_score_

best_params_rf, best_score_rf
print("Best Parameters for Random Forest:",
      grid_search_rf.best_params_)
print("Best Score:", grid_search_rf.best_score_)

Fitting 10 folds for each of 81 candidates, totalling 810 fits
Best Parameters for Random Forest: {'max_depth': 10,
'min_samples_leaf': 2, 'min_samples_split': 10, 'n_estimators': 50}
Best Score: 0.9022761760242792

```

The results from Grid Search for the Random Forest model show the following best hyperparameters:

max_depth: 10

min_samples_leaf: 2

min_samples_split: 10

n_estimators: 50

And the best score achieved with these parameters is approximately **90.23%**. This is a robust performance and indicates that the model is well-tuned with these parameters.

```

from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, classification_report

# Creating the Random Forest model with the best parameters from Grid Search
tuned_rf = RandomForestClassifier(n_estimators=50, max_depth=10,
min_samples_split=10,
                                min_samples_leaf=2, random_state=42)
tuned_rf.fit(X_train_no_id, y_train)

# Making predictions on the validation set
y_val_pred_tuned_rf = tuned_rf.predict(X_val_no_id)

# Evaluating the model
accuracy_tuned_rf = accuracy_score(y_val, y_val_pred_tuned_rf)
classification_report_tuned_rf = classification_report(y_val,
y_val_pred_tuned_rf)

```

```
# Print the results
print("Accuracy of Tuned Random Forest:", accuracy_tuned_rf)
print("\nClassification Report:\n", classification_report_tuned_rf)
```

Accuracy of Tuned Random Forest: 0.9023058252427184

Classification Report:

	precision	recall	f1-score	support
0	0.92	0.98	0.95	1466
1	0.63	0.28	0.39	182
accuracy			0.90	1648
macro avg	0.77	0.63	0.67	1648
weighted avg	0.88	0.90	0.89	1648

Hyperparameter Tuning

For hyperparameter tuning of the XGBoost model, we will use a similar approach as with the Random Forest model. However, XGBoost has its own set of key hyperparameters that can significantly influence its performance. We'll focus on a few important ones:

n_estimators: Number of gradient boosted trees. Increasing this number can make the model more complex and prone to overfitting.

max_depth: Maximum depth of each tree. Deeper trees can model more complex patterns but might lead to overfitting.

learning_rate: Step size shrinkage used in updating to prevent overfitting. Lower values make the boosting process more conservative.

subsample: Subsample ratio of the training instances. Setting it to 0.5 means that XGBoost would randomly sample half of the training data prior to growing trees.

```
from sklearn.model_selection import GridSearchCV
from xgboost import XGBClassifier

# Defining the parameter grid
param_grid_xgb = {
    'n_estimators': [50, 100, 150],
    'max_depth': [3, 5, 7],
    'learning_rate': [0.01, 0.1, 0.2],
    'subsample': [0.5, 0.7, 1.0]
}

# Initializing the Grid Search with XGBoost model
grid_search_xgb =
GridSearchCV(estimator=XGBClassifier(random_state=42),
```

```

param_grid=param_grid_xgb,
cv=10,
n_jobs=-1,
verbose=2)

# Fitting Grid Search to the data
grid_search_xgb.fit(X_train_no_id, y_train)

# Best parameters and best score
best_params_xgb = grid_search_xgb.best_params_
best_score_xgb = grid_search_xgb.best_score_

print("Best Parameters:", best_params_xgb)
print("Best Score:", best_score_xgb)

Fitting 10 folds for each of 81 candidates, totalling 810 fits
Best Parameters: {'learning_rate': 0.1, 'max_depth': 3,
'n_estimators': 100, 'subsample': 0.7}
Best Score: 0.9016691957511382

```

The results from Grid Search for the XGBoost model have identified the following optimal hyperparameters:

learning_rate: 0.1

max_depth: 3

n_estimators: 100

subsample: 0.7

With these parameters, the best score achieved is approximately **90.17%**. This is a strong performance and indicates that the model is well-tuned with these settings.

To evaluate the tuned XGBoost model, we'll follow similar steps as we did with the Random Forest model:

Create and Train the Tuned Model: We'll instantiate an XGBoost model using the optimal hyperparameters identified from Grid Search.

Make Predictions: Use the model to predict responses on the validation set. Evaluate Performance: We'll assess the model's performance using various metrics like accuracy, precision, recall, and the F1-score.

```

from xgboost import XGBClassifier
from sklearn.metrics import accuracy_score, classification_report

# Creating the XGBoost model with the best parameters from Grid Search
tuned_xgb = XGBClassifier(n_estimators=100, max_depth=3,
learning_rate=0.1, subsample=0.7, random_state=42)

```

```
tuned_xgb.fit(X_train_no_id, y_train)

# Making predictions on the validation set
y_val_pred_tuned_xgb = tuned_xgb.predict(X_val_no_id)

# Evaluating the model
accuracy_tuned_xgb = accuracy_score(y_val, y_val_pred_tuned_xgb)
classification_report_tuned_xgb = classification_report(y_val,
y_val_pred_tuned_xgb)

# Print the results
print("Accuracy of Tuned XGBoost:", accuracy_tuned_xgb)
print("\nClassification Report:\n", classification_report_tuned_xgb)
```

Accuracy of Tuned XGBoost: 0.9035194174757282

Classification Report:				
	precision	recall	f1-score	support
0	0.92	0.98	0.95	1466
1	0.63	0.30	0.41	182
accuracy			0.90	1648
macro avg	0.78	0.64	0.68	1648
weighted avg	0.89	0.90	0.89	1648

The evaluation results for the tuned XGBoost model show a promising performance:

Accuracy: Approximately 90.35%, which is quite good.

Precision and Recall for Class 1: The model has a precision of 0.63 and a recall of 0.30 for the minority class (class 1). This indicates that while the model is relatively precise in its predictions of class 1, it is only identifying 30% of the actual class 1 instances (low recall).

F1-Score for Class 1: The F1-score of 0.41 for class 1 reflects a balance between precision and recall but leans towards lower recall.

The XGBoost model for its higher overall accuracy, the next step is to apply this model to the test dataset to identify which potential customers should be targeted in the marketing campaign.

Steps: Load and Preprocess the Test Data: We need to ensure that the test data is preprocessed in the same way as the training data. This includes handling missing values, encoding categorical variables, and any other transformations that were applied.

Make Predictions: Use the tuned XGBoost model to predict on the test dataset.

Generate Output: We'll create a column in the test dataset with the predictions (1 for 'yes, market to this individual', 0 for 'no, do not market').

Consider Profit and Cost: When deciding whether to market to an individual, consider the cost of marketing (\$25 per customer) and the expected profit from customers who respond positively.

```
# Handling missing values: filling missing 'custAge' with median and
'schooling' with mode
test_df['custAge'].fillna(test_df['custAge'].median(), inplace=True)
test_df['schooling'].fillna(test_df['schooling'].mode()[0],
inplace=True)

# One-hot encoding for categorical variables
categorical_cols = ['profession', 'marital', 'schooling', 'default',
'housing', 'loan',
                    'contact', 'month', 'day_of_week', 'poutcome']
test_df_encoded = pd.get_dummies(test_df, columns=categorical_cols)

# Dropping the 'id' column as it was not used in training
test_df_encoded.drop(['id'], axis=1, inplace=True)

# Checking the first few rows after preprocessing
test_df_encoded.head()
```

	custAge	campaign	pdays	previous	emp.var.rate	cons.price.idx \
0	38.0	2	999	1	-1.1	94.199
1	35.0	2	3	1	-3.4	92.379
2	50.0	1	999	1	-1.8	92.893
3	30.0	1	999	0	1.4	93.444
4	39.0	1	999	0	-0.1	93.200

	cons.conf.idx	euribor3m	nr.employed	pmonths	...	month_oct
month_sep \						
0	-37.5	0.886	4963.6	999.0	...	0
1						
1	-29.8	0.788	5017.5	0.1	...	0
1						
2	-46.2	1.327	5099.1	999.0	...	0
0						
3	-36.1	4.964	5228.1	999.0	...	0
0						
4	-42.0	4.153	5195.8	999.0	...	0
0						

	day_of_week_fri	day_of_week_mon	day_of_week_thu	day_of_week_tue
\				
0	0	0	0	0
1	0	0	0	1
2	0	0	1	0

3	0	0	0	0
4	0	0	0	1

	day_of_week_wed	poutcome_failure	poutcome_nonexistent	poutcome_success
0	1	1	0	
0				
1	0	0	0	
1				
2	0	1	0	
0				
3	1	0	1	
0				
4	0	0	1	
0				

[5 rows x 64 columns]

To align the features of the test dataset with those of the training dataset, we need to perform the following steps:

Identify Missing Columns: Determine which columns are present in the training dataset but missing in the test dataset. Add these columns to the test dataset, filling them with zeros.

Remove Extra Columns: Identify any columns in the test dataset that are not present in the training dataset. Remove these columns.

Ensure Correct Order: Make sure the order of columns in the test dataset matches the order in the training dataset.

```
# List of columns from the training dataset
train_columns = list(X_train_no_id.columns)

# Add missing columns in test dataset
for col in train_columns:
    if col not in test_df_encoded.columns:
        test_df_encoded[col] = 0

# Ensure the order of columns in test dataset matches that of the
training dataset
test_df_encoded = test_df_encoded[train_columns]

# Remove columns in test dataset that are not in training dataset
test_df_encoded = test_df_encoded[train_columns]
```

Now that our test dataset is correctly aligned with our training dataset, we can proceed with making predictions using the tuned XGBoost model. This will allow us to identify which potential customers are most likely to respond positively to the marketing efforts.

```
# Making predictions on the test dataset
test_predictions = tuned_xgb.predict(test_df_encoded)
```

```
# Adding the predictions to the original test dataset
test_df['Marketing_Response'] = test_predictions
```

```
# Checking the first few rows with the predictions
test_df.head(100)
```

	custAge	profession	marital	schooling	default
housing loan \					
0	38.0	admin.	married	university.degree	no
no	yes				
1	35.0	services	married	high.school	no
no	no				
2	50.0	blue-collar	married	professional.course	unknown
yes	no				
3	30.0	admin.	single	university.degree	no
no	no				
4	39.0	services	divorced	high.school	no
yes	no				
..
..	...				
95	55.0	admin.	divorced	university.degree	no
no	no				
96	28.0	blue-collar	married	basic.9y	unknown
yes	no				
97	51.0	unemployed	divorced	university.degree	no
no	no				
98	49.0	management	married	university.degree	no
no	no				
99	31.0	unemployed	single	university.degree	no
yes	no				

	contact	month	day_of_week	...	poutcome	emp.var.rate	\
0	cellular	sep	wed	...	failure	-1.1	
1	cellular	sep	tue	...	success	-3.4	
2	cellular	may	thu	...	failure	-1.8	
3	cellular	aug	wed	...	nonexistent	1.4	
4	cellular	nov	tue	...	nonexistent	-0.1	
..	
95	telephone	may	tue	...	nonexistent	1.1	
96	cellular	jul	tue	...	nonexistent	1.4	
97	cellular	nov	thu	...	failure	-0.1	
98	cellular	nov	fri	...	nonexistent	-0.1	
99	cellular	nov	thu	...	nonexistent	-0.1	

	cons.price.idx	cons.conf.idx	euribor3m	nr.employed	pmonths
pastEmail \					
0	94.199	-37.5	0.886	4963.6	999.0

```

2
1          92.379          -29.8          0.788          5017.5          0.1
2
2          92.893          -46.2          1.327          5099.1          999.0
2
3          93.444          -36.1          4.964          5228.1          999.0
0
4          93.200          -42.0          4.153          5195.8          999.0
0
..          ...          ...          ...          ...          ...
...
95          93.994          -36.4          4.856          5191.0          999.0
0
96          93.918          -42.7          4.962          5228.1          999.0
0
97          93.200          -42.0          4.076          5195.8          999.0
1
98          93.200          -42.0          4.021          5195.8          999.0
0
99          93.200          -42.0          4.076          5195.8          999.0
0

   id  Marketing_Response
0    1                  0
1    2                  0
2    3                  0
3    4                  0
4    5                  0
..   ...                ...
95   96                  0
96   97                  0
97   98                  0
98   99                  0
99  100                  0

[100 rows x 23 columns]

```

Business Analysis

Count Positive Responses: Determine the number of customers predicted to respond positively (where Marketing_Response is 1). This will help in estimating the total number of customers we might want to target.

Calculate Total Profit: To calculate the total profit, use the formula:

Total Profit = (Number of Positive Responses × Average Profit per Responding Customer) - (Total Number of Customers Marketed × \$25)

For this calculation, you'll need the average profit per responding customer (avg_profit_per_respond). If you don't have this specific value, you can use a hypothetical or average value based on your business context.

Business Decision Making: Use these insights to make informed decisions about your marketing strategy. Depending on the estimated profit, decide whether it's cost-effective to market to all the customers predicted to respond positively.

```
# Counting the number of predicted positive responses
num_positive_responses = test_df['Marketing_Response'].sum()

# Assuming an average profit per positive response (you can replace
this with the actual value)
avg_profit_per_respond = 100 # Hypothetical value

# Calculating the total profit
total_profit = (num_positive_responses * avg_profit_per_respond) -
(len(test_df) * 25)

print("Number of Positive Responses:", num_positive_responses)
print("Total Profit:", total_profit)

Number of Positive Responses: 993
Total Profit: -724450
```

It looks like the calculation resulted in a significant negative total profit, which suggests that the cost of marketing to all the customers in this dataset outweighs the expected profit from positive responses.

Here are a few considerations and potential next steps:

Review Marketing Costs and Profit Margins: The calculated negative profit might indicate that the marketing costs are too high relative to the average profit per positive response. It's essential to review these figures to ensure they accurately reflect your business context.

Selective Marketing Strategy: Instead of marketing to all predicted positive responses, consider a more selective approach. You could target a subset of these customers, perhaps those with a higher probability of responding positively, to optimize the cost-effectiveness.

Reassess Model and Thresholds: The model's predictions and the threshold used for deciding a positive response (currently, a binary 0 or 1) could be re-evaluated. Adjusting the threshold to be more selective might help target customers more likely to respond.

Further Analysis of Predictions: Analyze the characteristics of customers predicted to respond positively. This might provide insights into which customer segments are more likely to be profitable, helping to refine the targeting strategy.

Business Strategy Review: It may be necessary to review the overall marketing and customer acquisition strategy, considering the balance between costs and potential returns.

The negative profit could also result from the hypothetical average profit per responding customer. If this value isn't accurate, adjusting it to a more realistic figure could significantly change the profit calculation.

```
# Summary details
num_positive_responses = 993 # Number of positive responses
total_customers_marketed = 32950 # Total number of customers marketed
to
total_profit = -674800 # Total profit (or loss)

# Creating the summary text
summary_text = f"""
Marketing Campaign Analysis Summary
-----
Total Customers Marketed To: {total_customers_marketed}
Number of Predicted Positive Responses: {num_positive_responses}
Average Profit Per Responding Customer: $150

Total Profit (or Loss): ${total_profit}

Note: The total profit calculation is based on the assumption of an
average profit of $150 per responding customer
and a marketing cost of $25 per customer.
"""

# Printing the summary
print(summary_text)

Marketing Campaign Analysis Summary
-----
Total Customers Marketed To: 32950
Number of Predicted Positive Responses: 993
Average Profit Per Responding Customer: $150

Total Profit (or Loss): $-674800

Note: The total profit calculation is based on the assumption of an
average profit of $150 per responding customer
and a marketing cost of $25 per customer.
```

This result suggests that the marketing strategy, as modeled, would not be profitable under these conditions.